

# Software Systems Research – Portfolio Review

Dr. Nelson R. Manohar Alers

[nelsonmanohar@yahoo.com](mailto:nelsonmanohar@yahoo.com)

[nelsonmanohar.sphosting.com/research](http://nelsonmanohar.sphosting.com/research)



# Outline of the Talk – Update

- Background [10%]
- Computer-Supported Collaboration [25%]
- Dynamically Customized Web Touring [25%]
- **Multimedia Computing Networking [35%]**
  - Building Robust Network Performance Indicator
  - **Distributed Resource Management for Multimedia**
- Wrap-Up [5%]



# Distributed Resource Management

Work at IBM T. J. Watson Research Center

By N. R. Manohar, L. Lumelsky, and S. Wood  
U.S. Patents: 6,516,350; 6,463,454; 6,529,950;  
6,377,996; 6,460,082; 6,466,980;



# Distributed Resource Management Plane



## ■ From: Ad-Hoc Resource Management

- ad-hoc client-services
  - streaming, metering, etc.
- ad-hoc server-services
  - caching, load-balancing, etc.

## ■ To: Resource Management Plane

- distributed management plane
- brokering of clients to servers
- standard plane server-services
  - caching, replication, load-balancing, dynamic hosting, etc.
- compliant plug-in servers

**FIG. 3**





# Motivation and Goals

## ■ Evaluation

- Multimedia Computing Networking at Internet2-like networking level
  - very high performance Backbone (vBNS 2.4Gbps OC-12)
  - differentiated class services, RSVP, RTCP, RTP, etc.

## ■ What could benefit from this?

- large multimedia objects under a value-asset model
- e.g., movies

## ■ What could we do now?

- distributed resource management



# Very Large Multimedia Objects?

## ■ internet-perspective

- flow-oriented, end-to-end, distributed QoS/reservation problem
- non-negligible, considerable, replication cost

## ■ web-perspective

- extremely large in comparison to dominant web objects
- request characterization of web-servers [MSNBC00]
  - some objects significantly more requested than others
  - small set of objects accounts for majority of requests (Zipf relative frequencies)
- spatial and temporal locality on requests [IBM97]
  - movie is blockbuster, demographic and temporal consequences
- and then, some are exceptions, i.e., statistical outliers
  - for example, fads (30s movies, some-actor movies, etc.)

## ■ provisioning-perspective

- value-asset model, digital rights model
- non-cacheable

# Distributed Resource Management?


- What else can we do there? – again, viewpoint matters
- adapting demand to the network (variable client demand, fixed server capacities)
  - balance client demand presented to the network *wrt* fixed server capacities found across the network
  - traditional load balancing, i.e., demand shaping
  - that is, “demand-follows-capacity” policy
- **adapting the network to demand (variable client demand, variable server capacity)**
  - regulate the allocation of server capacities across the network *wrt* client demand presented to the network
  - the other side of the coin, i.e., distributed capacity shaping
  - that is, “capacity-follows-demand” policy





# Capacity Shaping – Basic Idea


- **capacity shaping mechanism**
  - object is associated with object replicas
  - object replica specifies (serving) resource requirements (i.e., normalized capacity)
  - allocation of an object replica commits the replica's resource allocation (i.e., normalized capacity allocation)
- **distributed dynamic capacity shaping**
  - expiration time associated with object replica
  - server placement of object replicas varies over time
  - total number of object replicas varies over time
- **thus, allocated object-serving capacity placed over the network shaped over time**
  - done for selected objects – referred to as hot-objects





# Self-Regulated Capacity Shaping

- load balancing and capacity shaping are complementary
- traditionally, throttle control (against fixed capacity) for demand-shaping
- now, self-regulation (against variable capacity) for capacity shaping
- building block: robust state signaling




- dynamic capacity shaping for which objects?
  - Zipf relative frequency distribution – hot objects
- then, how is self-regulated capacity shaping implemented?





# Capacity Shaping Implementation

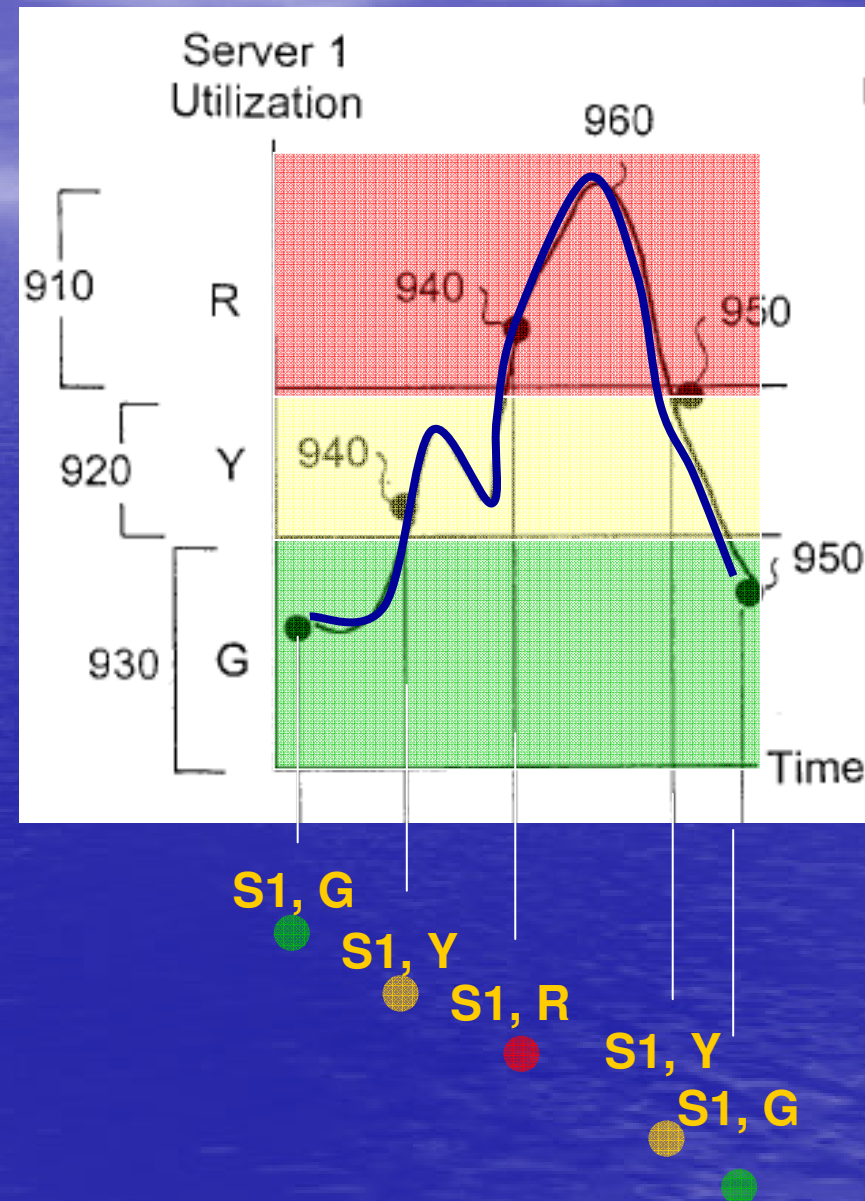
1. scarcity check
  - demand/capacity analysis (a)
2. replica creation
3. replica placement
4. replication
5. oversupply check
  - demand/capacity analysis (b)
6. replica deletion





# Management of Remote Capacities


- low overhead on tracking of distributed state
  - server capacity is normalized
  - semaphore-like
  - green, red, yellow regions
  - trigger-management problem
- then, to handle transient changes
  - regions act as buffering regions
  - green
    - safe operating region
    - low management overhead region
  - red
    - critical spare ( $X\%$ ) region
    - non-allocable (safe operating margin)
  - yellow
    - trigger region and buffering region
    - buffers transient state changes and messaging delays





# Management of Heterogeneous Servers

- to handle server heterogeneity
  - regions are server specified
- management goal becomes about
  - handling of red servers
- where
  - forewarning through state signaling
  - capacity spare used to manage signaling



# Resource Management at Plane

## ■ Low-overhead State Tracking

- quantized demand
- normalized state & capacity

## ■ Object to Demand

- demand volume
- object-demand rating
- interval-tracking

## ■ Object to Replica

- server placement
- server capacity
- expiration time

## ■ Server to Capacity

- capacity rating
- utilization state
- globality

ObjectID	Demand rate req/s	Volume $t_{i-1}$ req	Volume $t_i$ req	Hot Object	Time Stamp
420	10	120	60	yes	$t_1$
425	5	60	55	no	$t_1$
428	5	30	62	no	$t_0$

object 420: hot, large demand

Object_ID	Replica	Server	Transient Replica	Time-to-Live
420	421	1211	NO	
	422	1221	YES	060599-133000
440	441	1211	NO	

transient replica for 420 at server 1221

Server	IP Address	Capacity Rating	Utilization State	Timestamp	Globality
1211	209.09.9.127	Low	Red	$t_1$	local
1221	128.0.0.1	High	Green	$t_2$	global


global server 1221: high capacity, low use





# Resource Management at Server


- **reservation bins abstraction**
  - normalized management of replication and provisioning problems
  - storage-bins – for replica placement problem
  - service-bin – for resource reservation problem
- **reservation pools abstraction**
  - for local resource management
    - local serving resources
  - for global resource management
    - *degree-of-freedom* resource for the distributed capacity-shaping problem



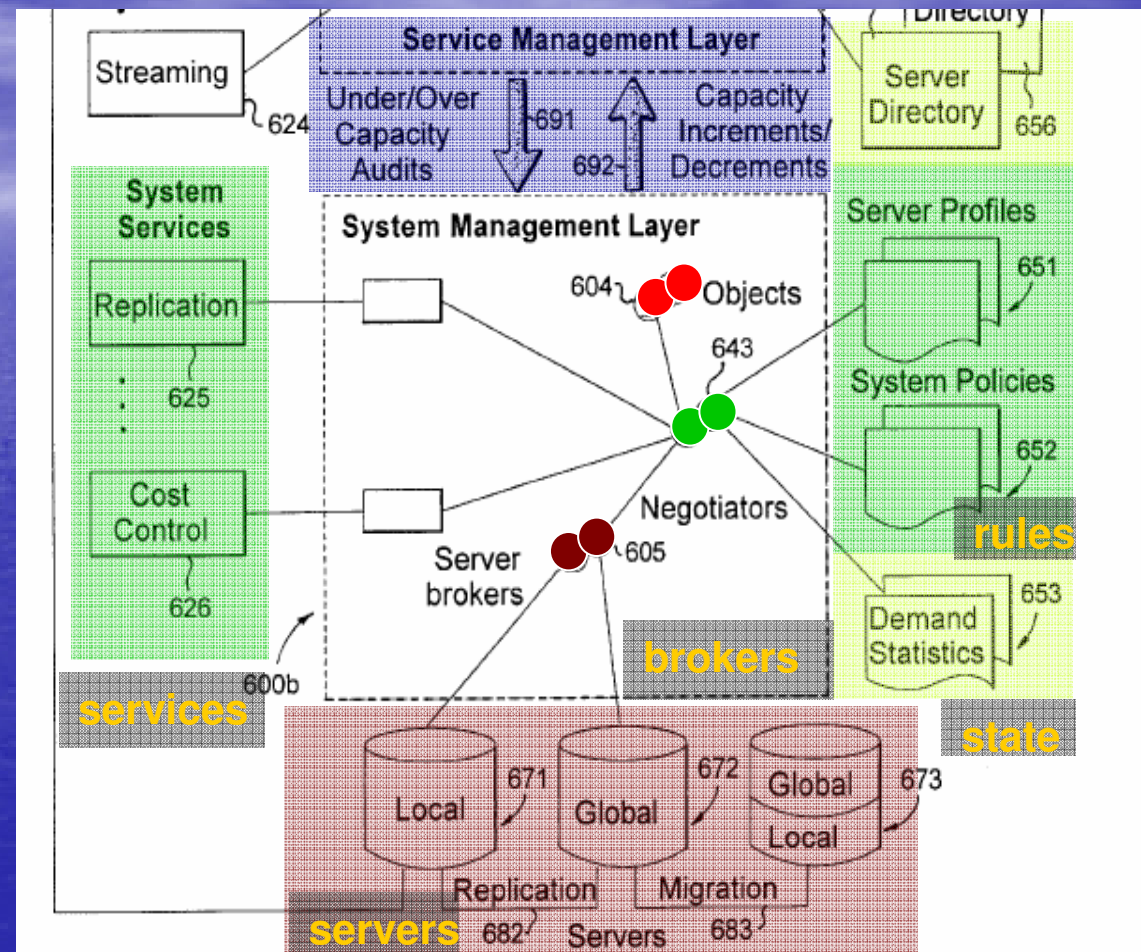


# Capacity Shaping Management Layer

- capacity shaping plane brokers object replication between servers
  - dynamically creating capacity for load balancing

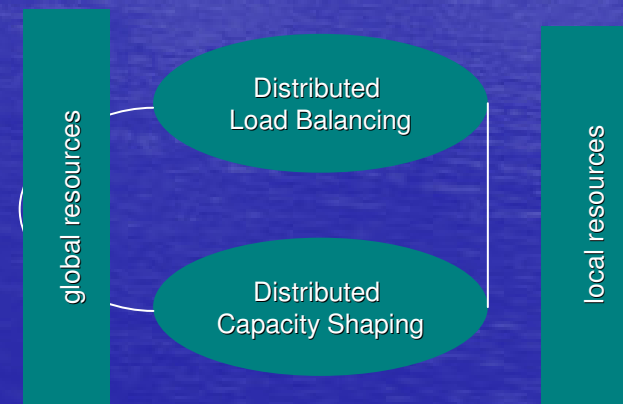


- driven by service management layer
  - by demand-capacity analysis resulting in dynamic replica hosting at brokered servers



# Contributions

- **Distributed Dynamic Capacity Shaping**
  - for distributed resource management of large (multimedia) objects over high-bandwidth internet
  - self-regulated distributed capacity shaping
    - integrated load-balancing and dynamic capacity shaping
    - management of global and local replicas *wrt* server capacities
  - distributed state management
    - red-server avoidance management
    - low-overhead distributed state tracking
    - stability-oriented trigger signaling
    - tradeoff of local capacity to remote state tracking
- **complementary to existing technologies**
  - multicast, rights, RSVP, grids, etc.
  - but leveraging high bandwidth internet (e.g., internet2)



# Related Work

- **Dynamic Capacity Virtual Hosting**
  - server farms (e.g., IBM Global Services) – demand-shaping of capacity at centralized center(s) vs. capacity-shaping of demand through placement of distributed capacities
- **Distributed Programmable Planes**
  - grid computing (e.g., Globus) – resource management for entirely different problem, small number of large (computations and datasets) objects with low viewership vs. large number of large objects with high and wide viewership
- **Distributed Replica Management**
  - edge-caching network (e.g., Akamai) – on-demand caching of relatively small objects vs. valuable asset very large object
- **Multicast Streaming**
  - multicast is complementary technology – plane represents the self-regulated placement of the sources of the multicast trees
- **Brokered Distributed QoS Architectures**
  - distributed QoS brokering is complementary technology – delivering end-to-end QoS feasibility used for brokering of clients to servers

